

Dang Nguyen

Email: dangnguyen@uchicago.edu | Website: dangng2004.github.io/ |  [Google Scholar](#) | Phone: +1 (773) 936-0857

Research Interests

- **AI values:** study the judgment and preferences of AI models and how they can influence humans’ digital norms.
- **Alignment:** create methods to align models to human values to ensure AI safety.
- **Human-centered AI:** understand human-AI interaction and design AI that maximizes human potential.

Education

2024 - Exp. 2029	Computer Science Ph.D. , The University of Chicago. Advisor: Prof. Chenhao Tan
2020 - 2024	Computer Science B.S. & Mathematics B.A. , The University of Chicago. Specialization in Machine Learning.

Preprints

2026	Benchmarking Agentic Review Systems. [Link] Nguyen, D., Hao, W., Elazar, Y., & Tan, C. <i>Oral presentation (top 5%) at the AI-Native Academia Workshop at NeurIPS 2026.</i>
2026	Moral Mazes in the Era of Large Language Models. [Link] Nguyen, D., Fu, H.Y., West, P., Holtzman, A., & Tan, C. <i>Accepted to the P-AGI Workshop at ICLR 2026.</i>
2026	Iterative Finetuning is Mostly Idempotent. [Link] Roe, Z.*, Sanderson, J.*, Nguyen, D.*, Huang, J.*, Nief, T., Shrivastava, A., Tan, C., & Holtzman, A. <i>Under review.</i>

Selected Publications

2025	On the Effectiveness and Generalization of Race Representations for Debiasing High-Stakes Decisions. [Link] Nguyen, D., & Tan, C. In Conference on Language Modeling (COLM). <i>Top 1% of papers in the review process. Best Poster Award runner-up at Midwest Machine Learning Symposium 2025.</i> Keywords: fairness and bias, mechanistic interpretability, causal interventions
2025	GPT-4V Cannot Generate Radiology Reports Yet. [Link] Jiang, Y.*, Chen, C.*, Nguyen, D., Mervak, B., & Tan, C. In Findings of the Association for Computational Linguistics (NAACL). <i>Lightning talk at the Toyota Technological Institute at Chicago (TTIC) Multimodal AI Workshop 2024.</i> Keywords: evaluation of large vision-language models, error analysis
2023	Pragmatic Radiology Report Generation. [Link] Nguyen, D., Chen, C., He, H., & Tan, C. In Proceedings of Machine Learning for Health (ML4H). <i>Example paper for ML4H 2024.</i> Keywords: vision-language modeling, finetuning LLMs, reducing hallucinations

Research Experience

2026 - Present	OpenAIReview: An Open-sourced AI Reviewer. Advisor: Prof. Chenhao Tan • Benchmarking Agentic Review Systems.
2025 - Present	Communication and Intelligence , Graduate Researcher. Advisor: Prof. Ari Holtzman, Prof. Chenhao Tan • Moral Mazes in the Era of Large Language Models. • Iterative Finetuning is Mostly Idempotent.

2024 - 2025	Chicago Human+AI Lab , Graduate Researcher. Advisor: Prof. Chenhao Tan On the Effectiveness and Generalization of Race Representations for Debiasing High-Stakes Decisions.
2023 - 2024	Chicago Human+AI Lab , Undergraduate Researcher. Advisor: Prof. Chenhao Tan Pragmatic Radiology Report Generation.
2020 - 2022	Machine Learning for Systems Research (UCARE Lab) , Undergraduate Researcher. Advisor: Prof. Haryadi Gunawi Developed statistical and ML-based methods (based on K-means, DBSCAN, KL divergence) for detecting anomalous hard disk drives at scale.

Industry Experience

Summer 2026	Atoms , Machine Learning PhD Intern. Developed an LLM-assisted framework for non-experts to achieve expert-level annotation of complex conversation data.
Summer 2022	NetApp Inc. , Research Intern. Developed ML algorithms for slow disk detection, improving the true positive rate from 10% to 59% over the existing threshold-based detector.

Teaching Experience

Spring 2026	Teaching Assistant , CMSC 25750: Large Language Models, University of Chicago. Designed homework assignments and mentored 25 student teams on final projects.
Winter 2026	Teaching Assistant , CMSC 25700: Natural Language Processing, University of Chicago. Designed homework assignments and mentored 9 student teams on final projects.
Winter 2025	Teaching Assistant , CMSC 25700: Natural Language Processing, University of Chicago. Designed homework assignments and mentored 7 student teams on final projects. 3 advised teams ranked in the top 6 out of 26; one won the Best Final Report award.

Service

2026	Reviewer , Conference on Language Modeling (COLM).
2026	Reviewer , Association for Computational Linguistics (ACL) System Demonstrations.
2025	Reviewer , Conference on Language Modeling (COLM).
2025	Reviewer , International Conference on Learning Representations (ICLR).
2024	Reviewer , Machine Learning for Health (ML4H).